

Intruder Detection - Distributed Systems

Alan Gallo, Cedric Sillaber, Frantisek Sova

January 2025

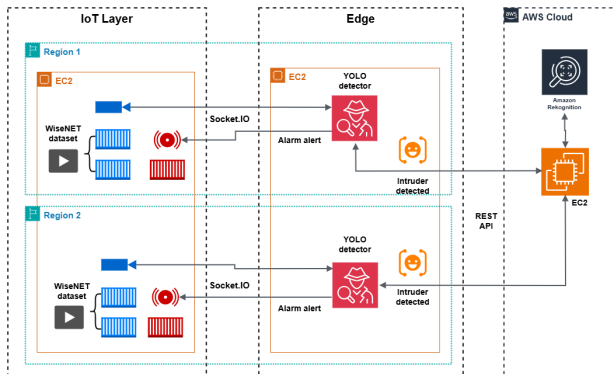


Content

- System architecture
- Implementation details
- Demo
- Evaluation

System architecture

Final system architecture:



IoT Implementation Details

Video files stored on iot, opened with opencv

Frame Processing

- extract 1 frame per second from 30 fps video stream
- main bottle neck!

Communication using Socket.IO

- websockets, built ontop of application layer
- Good for continuous dataflow, persistent connection
- Setup once, then just send data



socket.io

Edge Implementation Details

Async Processing Pipeline

- two workers working asynchronously:
 - worker 1: frame buffer queue
 - worker 2: process buffer: YOLO person detection, cloud communication

Communication: Flask REST

- REST API for communication between edge and cloud
- Edge sends HTTP request to Cloud → Cloud sends HTTP response

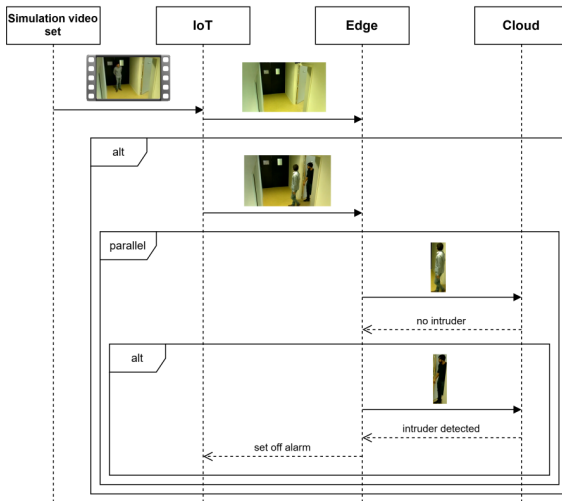


Cloud Implementation Details

- REST API endpoint
- AWS Rekognition



Example Workflow/ Controlgraph



Evaluation


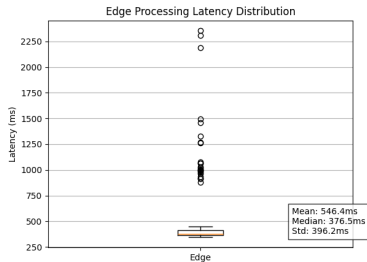
Experiment Configuration:

- *set3* from WiseNET dataset
- 2 edge devices
- 2 cams + alarm per edge
- cloud instance

Metrics:


- Edge processing latency (yolo processing)
- Cloud latency (AWS Rekognition)
- round-time trip time from camera to alarm

Evaluation - Edge & Cloud Latencies




Evaluation - Round Trip Time

Round trip time: camera frame to alarm.




Evaluation - System Bottleneck

- System bottleneck on edge
- new frames / second $>$ frame processing rate \rightarrow very high latency



RTT for 2 frames per second
sent



RTT for 0.5 frames per second
sent