

# Intruder Detection - Distributed Systems

Alan Gallo, Cedric Sillaber, Frantisek Sova

January 2025

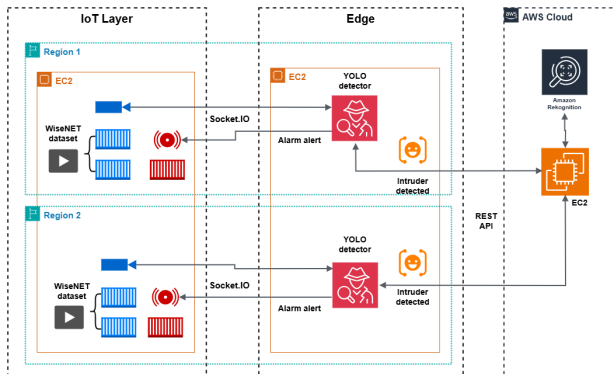


# Content

- System architecture
- Implementation details
- Demo
- Evaluation

# System architecture

Final system architecture:



# IoT Implementation Details

Video files stored on iot, opened with opencv

## Frame Processing

- extract 1 frame per second from 30 fps video stream
- main bottle neck!

## Communication using Socket.IO

- websockets, built ontop of application layer
- Good for continuous dataflow, persistent connection
- Setup once, then just send data



socket.io

# Edge Implementation Details

## Async Processing Pipeline

- two workers working asynchronously:
  - worker 1: frame buffer queue
  - worker 2: process buffer: YOLO person detection, cloud communication

## Communication: Flask REST

- REST API for communication between edge and cloud
- Edge sends HTTP request to Cloud → Cloud sends HTTP response

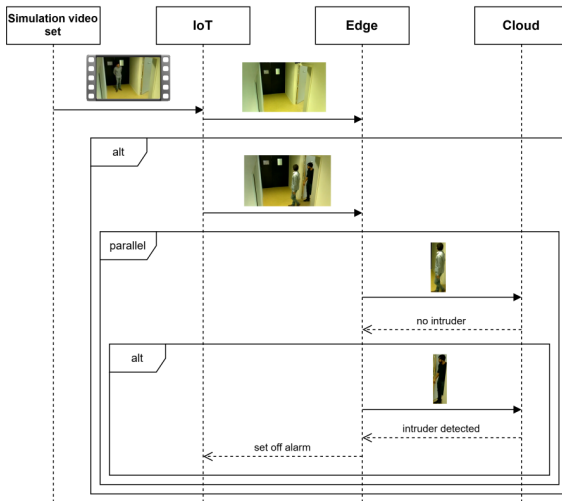


# Cloud Implementation Details

- REST API endpoint
- Wrapper for AWS Rekognition



# Example Workflow/ Controlgraph



# Evaluation

## Experiment Configuration:

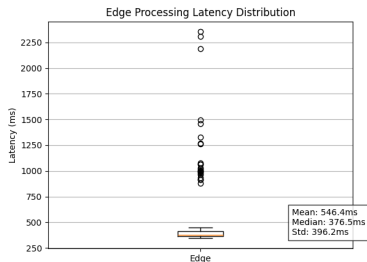
- *set3* from WiseNET dataset
- 2 edge devices
- 2 cams + alarm per edge
- cloud instance

## Metrics:


- Edge processing latency (yolo processing)
- Cloud latency (AWS Rekognition)
- Round-trip time from camera to alarm



# Evaluation - Edge & Cloud Latencies




YOLO person detection latency



Cloud latency (AWS  
Rekognition)


# Evaluation - Round Trip Time

Round-trip time: camera frame to alarm.




# Evaluation - System Bottleneck

- System bottleneck on edge
- New frames / second  $>$  frame processing rate  $\rightarrow$  very high latency



RTT for 2 frames per second  
sent



RTT for 0.5 frames per second  
sent